

# Applying Augmentation, Fine-Tuning Last Dense Block, Transition Layer, Classifier

In this notebook, we will only apply **augmentation**, **fine-tune the last dense block**, **transition layer**, and **classifier**. We'll use this approach as a comparison with other notebooks that use different techniques.

```
In [1]: import os
import copy
import pandas as pd
import torch
import torchvision
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
import torchvision.transforms as transforms
import torchvision.models as models
from PIL import Image
from sklearn.preprocessing import StandardScaler, LabelEncoder
import numpy as np
from sklearn.metrics import roc_auc_score
import torch.nn.functional as F
```

```
In [2]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device.type}")
```

Using device: cuda

```
In [3]: CKPT_PATH = '../model.pth.tar'
N_CLASSES = 14
CLASS_NAMES = [ 'Atelectasis', 'Cardiomegaly', 'Effusion', 'Infiltration', 'Mass',
                'Pneumothorax', 'Consolidation', 'Edema', 'Emphysema', 'Fibrosis',
DATA_DIR = '../images'
TEST_IMAGE_LIST = '../test_list.txt'
BATCH_SIZE = 64
```

```
In [4]: class ChestXrayDataSet(Dataset):
    def __init__(self, data_dir, image_list_file, transform=None):
        """
        Args:
            data_dir: path to image directory.
            image_list_file: path to the file containing images
                        with corresponding labels.
            transform: optional transform to be applied on a sample.
        """
        image_names = []
        labels = []
        with open(image_list_file, "r") as f:
            for line in f:
                items = line.split()
                image_name = items[0]
                label = items[1:]
```

```

        label = [int(i) for i in label]
        image_name = os.path.join(data_dir, image_name)
        image_names.append(image_name)
        labels.append(label)

    self.image_names = image_names
    self.labels = labels
    self.transform = transform

    def __getitem__(self, index):
        """
        Args:
            index: the index of item

        Returns:
            image and its labels
        """
        image_name = self.image_names[index]
        image = Image.open(image_name).convert('RGB')
        label = self.labels[index]
        if self.transform is not None:
            image = self.transform(image)
        return image, torch.FloatTensor(label)

    def __len__(self):
        return len(self.image_names)

```

```

In [5]: class DenseNet121(nn.Module):
        """Model modified.

        The architecture of our model is the same as standard DenseNet121
        except the classifier layer which has an additional sigmoid function.

        """
        def __init__(self, out_size):
            super(DenseNet121, self).__init__()
            self.densenet121 = torchvision.models.densenet121(pretrained=True)
            num_fts = self.densenet121.classifier.in_features
            self.densenet121.classifier = nn.Sequential(
                nn.Linear(num_fts, out_size),
                nn.Sigmoid()
            )

        def forward(self, x):
            x = self.densenet121(x)
            return x

```

```

In [6]: def compute_AUCs(gt, pred):
        """Computes Area Under the Curve (AUC) from prediction scores.

        Args:
            gt: Pytorch tensor on GPU, shape = [n_samples, n_classes]
                true binary labels.
            pred: Pytorch tensor on GPU, shape = [n_samples, n_classes]
                can either be probability estimates of the positive class,

```

```

        confidence values, or binary decisions.

Returns:
    List of AUROCs of all classes.
"""
AUROCs = []
gt_np = gt.cpu().numpy()
pred_np = pred.cpu().numpy()
for i in range(N_CLASSES):
    AUROCs.append(roc_auc_score(gt_np[:, i], pred_np[:, i]))
return AUROCs

```

```

In [7]: import torchvision.transforms as transforms

TRAIN_LIST = "../train_list.txt"
VALID_LIST = "../val_list.txt"
IMAGE_DIR = "../images"

data_transforms = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std=[0.229, 0.224, 0.225])
])

train_dataset = ChestXrayDataSet(IMAGE_DIR, TRAIN_LIST, transform=data_transforms)
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=64, shuffle=False)

```

```

In [8]: images, labels = next(iter(trainloader))
print("Images shape:", images.shape)
print("Labels shape:", labels.shape)

```

```

Images shape: torch.Size([64, 3, 224, 224])
Labels shape: torch.Size([64, 14])

```

```

In [9]: model = DenseNet121(N_CLASSES).to(device)
model = torch.nn.DataParallel(model).to(device)

if os.path.isfile(CKPT_PATH):
    print("> loading checkpoint")
    modelCheckpoint = torch.load(CKPT_PATH)['state_dict']
    for k in list(modelCheckpoint.keys()):
        index = k.rindex('.')
        if (k[index - 1] == '1' or k[index - 1] == '2'):
            modelCheckpoint[k[:index - 2] + k[index - 1:]] = modelCheckpoint[k]
            del modelCheckpoint[k]
    model.load_state_dict(modelCheckpoint)
    print("> loaded checkpoint")
else:
    print("> no checkpoint found")

normalize = transforms.Normalize([0.485, 0.456, 0.406],
                                 [0.229, 0.224, 0.225])

test_dataset = ChestXrayDataSet(data_dir=DATA_DIR,

```

```

        image_list_file=TEST_IMAGE_LIST,
        transform=transforms.Compose([
            transforms.Resize(256),
            transforms.TenCrop(224),
            transforms.Lambda
                (lambda crops: torch.stack([transforms.ToTensor(crop)
                    for crop in crops])),
            transforms.Lambda
                (lambda crops: torch.stack([normalize(crop)
                    for crop in crops]))
        ])
test_loader = DataLoader(dataset=test_dataset, batch_size=BATCH_SIZE,
                        shuffle=False, num_workers=0, pin_memory=True)

```

c:\Users\aroc\miniforge3\envs\572\Lib\site-packages\torchvision\models\\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

warnings.warn(c:\Users\aroc\miniforge3\envs\572\Lib\site-packages\torchvision\models\\_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=DenseNet121\_Weights.IMAGENET1K\_V1`. You can also use `weights=DenseNet121\_Weights.DEFAULT` to get the most up-to-date weights.

warnings.warn(msg)  
=> loading checkpoint  
=> loaded checkpoint

```

In [10]: # Data transformations
train_transforms = transforms.Compose([
    transforms.Resize(256),
    transforms.RandomResizedCrop(224),
    transforms.RandomRotation(10),
    transforms.ColorJitter(brightness=0.2, contrast=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

valid_transforms = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

```

```

In [11]: # Class counts and weights
class_counts = [
    959, # Atelectasis
    276, # Cardiomegaly
    972, # Effusion
    1591, # Infiltration
    313, # Mass
    461, # Nodule
    140, # Pneumonia
    469, # Pneumothorax
    397, # Consolidation
    140, # Edema
    208, # Emphysema
    244, # Fibrosis
]

```

```

    349, # Pleural_Thickening
    35  # Hernia
]

total_count = sum(class_counts)
class_weights = [total_count / count for count in class_counts]
class_weights_cpu = torch.tensor(class_weights, dtype=torch.float)

# Compute sample weights
sample_weights = []
for _, label in train_dataset:
    label = label.float()
    weight = torch.sum(class_weights_cpu * label).item()
    sample_weights.append(weight)

sample_weights = torch.tensor(sample_weights, dtype=torch.float)
sampler = WeightedRandomSampler(weights=sample_weights, num_samples=len(sample_weights))
class_weights_gpu = torch.tensor(class_weights, dtype=torch.float).to(device)

```

In [12]: *# Replace the classifier with a more regularized version*

```

old_classifier = model.module.densenet121.classifier
num_fts = None
if isinstance(old_classifier, nn.Sequential):
    # Get input features from the first Linear layer
    for layer in old_classifier:
        if isinstance(layer, nn.Linear):
            num_fts = layer.in_features
            break
else:
    # Direct Linear layer
    num_fts = old_classifier.in_features

model.module.densenet121.classifier = nn.Sequential(
    nn.Linear(num_fts, 1024),
    nn.ReLU(),
    nn.Linear(1024, 512),
    nn.ReLU(),
    nn.Dropout(0.3),
    nn.Linear(512, 256),
    nn.ReLU(),
    nn.Dropout(0.2),
    nn.Linear(256, N_CLASSES),
    nn.Sigmoid()
)

print("Replaced classifier with a regularized version")

```

Replaced classifier with a regularized version

In [13]: *# Freeze all parameters*

```

for param in model.module.densenet121.parameters():
    param.requires_grad = False

# Unfreeze the classifier
for param in model.module.densenet121.classifier.parameters():
    param.requires_grad = True

```

```

# Also unfreeze the last dense block (denseblock4)
for param in model.module.densenet121.features.denseblock4.parameters():
    param.requires_grad = True

# Also unfreeze the transition layer before denseblock4
for param in model.module.densenet121.features.transition3.parameters():
    param.requires_grad = True

# Count trainable parameters
trainable_params = [p for p in model.parameters() if p.requires_grad]
print(f"Number of trainable parameters: {sum(p.numel() for p in trainable_params):,},

```

Number of trainable parameters: 4,393,742

```

In [14]: criterion = nn.BCELoss()

# Add weight decay to combat overfitting
optimizer = optim.Adam(
    filter(lambda p: p.requires_grad, model.parameters()),
    lr=1e-4,
    weight_decay=1e-5
)

# Prepare datasets
train_dataset = ChestXrayDataSet(IMAGE_DIR, TRAIN_LIST, transform=train_transforms)
valid_dataset = ChestXrayDataSet(IMAGE_DIR, VALID_LIST, transform=valid_transforms)

# Create data Loaders
trainloader = DataLoader(train_dataset, sampler=sampler, batch_size=64, pin_memory=True)
validloader = DataLoader(valid_dataset, batch_size=64, shuffle=False, pin_memory=True)

```

```

In [15]: def trainer(model, criterion, optimizer, trainloader, validloader, epochs=20, patience=10):
    train_loss, valid_loss, valid_accuracy = [], [], []

    # To track best model
    best_aucroc = 0.0
    best_model_weights = None
    counter = 0 # For early stopping

    for epoch in range(epochs):
        # Training Phase
        model.train()
        epoch_train_loss = 0.0
        for inputs, labels in trainloader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()

            outputs = model(inputs)
            loss = criterion(outputs, labels)

            loss.backward()
            optimizer.step()
            epoch_train_loss += loss.item() * inputs.size(0)

        train_loss.append(epoch_train_loss / len(trainloader.dataset))

```

```

# Validation Phase
model.eval()
epoch_valid_loss = 0.0
all_labels = torch.FloatTensor().to(device)
all_outputs = torch.FloatTensor().to(device)

with torch.no_grad():
    for inputs, labels in validloader:
        inputs, labels = inputs.to(device), labels.to(device)

        outputs = model(inputs)
        loss = criterion(outputs, labels)

        epoch_valid_loss += loss.item() * inputs.size(0)

        all_labels = torch.cat((all_labels, labels), 0)
        all_outputs = torch.cat((all_outputs, outputs), 0)

# Calculate metrics
predictions = (all_outputs > 0.5).float()
correct = (predictions == all_labels).sum().item()
total = all_labels.numel()
accuracy = correct / total

aurocs = compute_AUCs(all_labels, all_outputs)
mean_auroc = np.mean(aurocs)

valid_loss.append(epoch_valid_loss / len(validloader.dataset))
valid_accuracy.append(accuracy)

if verbose:
    print(f"Epoch {epoch+1}/{epochs} - Train Loss: {train_loss[-1]:.4f}, Val Loss: {valid_loss[-1]:.4f}")
    print(f"Accuracy: {accuracy:.4f}, Mean AUROC: {mean_auroc:.4f}")

    if (epoch + 1) % 5 == 0 or epoch == epochs - 1:
        for i, auroc in enumerate(aurocs):
            print(f" {CLASS_NAMES[i]}: AUROC = {auroc:.4f}")

# Early stopping and model saving
if mean_auroc > best_auroc:
    best_auroc = mean_auroc
    best_model_weights = copy.deepcopy(model.state_dict())
    print(f" New best model with AUROC: {best_auroc:.4f}")
    counter = 0 # Reset counter
else:
    counter += 1
    if counter >= patience:
        print(f" Early stopping triggered after {epoch+1} epochs")
        break

# Load the best model weights
if best_model_weights is not None:
    model.load_state_dict(best_model_weights)
    print(f"Loaded best model with AUROC: {best_auroc:.4f}")

```

```
return {  
    'train_loss': train_loss,  
    'valid_loss': valid_loss,  
    'valid_accuracy': valid_accuracy,  
    'best_auroc': best_auroc  
}
```

```
In [16]: model.to(device)  
trainer(model, criterion, optimizer, trainloader, validloader, epochs=20)
```



Epoch 1/20 - Train Loss: 0.5128, Valid Loss: 0.3110  
Accuracy: 0.8932, Mean AUROC: 0.5464  
New best model with AUROC: 0.5464  
Epoch 2/20 - Train Loss: 0.3846, Valid Loss: 0.3070  
Accuracy: 0.8932, Mean AUROC: 0.6059  
New best model with AUROC: 0.6059  
Epoch 3/20 - Train Loss: 0.3779, Valid Loss: 0.2933  
Accuracy: 0.8932, Mean AUROC: 0.6986  
New best model with AUROC: 0.6986  
Epoch 4/20 - Train Loss: 0.3703, Valid Loss: 0.2840  
Accuracy: 0.8960, Mean AUROC: 0.7465  
New best model with AUROC: 0.7465  
Epoch 5/20 - Train Loss: 0.3605, Valid Loss: 0.2764  
Accuracy: 0.8963, Mean AUROC: 0.7712  
Atelectasis: AUROC = 0.7912  
Cardiomegaly: AUROC = 0.9431  
Effusion: AUROC = 0.8185  
Infiltration: AUROC = 0.6108  
Mass: AUROC = 0.8355  
Nodule: AUROC = 0.7722  
Pneumonia: AUROC = 0.5734  
Pneumothorax: AUROC = 0.8203  
Consolidation: AUROC = 0.7534  
Edema: AUROC = 0.8917  
Emphysema: AUROC = 0.7355  
Fibrosis: AUROC = 0.6895  
Pleural\_Thickening: AUROC = 0.7631  
Hernia: AUROC = 0.7987  
New best model with AUROC: 0.7712  
Epoch 6/20 - Train Loss: 0.3551, Valid Loss: 0.2689  
Accuracy: 0.8980, Mean AUROC: 0.7917  
New best model with AUROC: 0.7917  
Epoch 7/20 - Train Loss: 0.3526, Valid Loss: 0.2626  
Accuracy: 0.8962, Mean AUROC: 0.8019  
New best model with AUROC: 0.8019  
Epoch 8/20 - Train Loss: 0.3487, Valid Loss: 0.2585  
Accuracy: 0.8996, Mean AUROC: 0.8058  
New best model with AUROC: 0.8058  
Epoch 9/20 - Train Loss: 0.3423, Valid Loss: 0.2542  
Accuracy: 0.8988, Mean AUROC: 0.8124  
New best model with AUROC: 0.8124  
Epoch 10/20 - Train Loss: 0.3407, Valid Loss: 0.2562  
Accuracy: 0.8969, Mean AUROC: 0.8142  
Atelectasis: AUROC = 0.8301  
Cardiomegaly: AUROC = 0.9437  
Effusion: AUROC = 0.8356  
Infiltration: AUROC = 0.6672  
Mass: AUROC = 0.8776  
Nodule: AUROC = 0.8332  
Pneumonia: AUROC = 0.5787  
Pneumothorax: AUROC = 0.8509  
Consolidation: AUROC = 0.7767  
Edema: AUROC = 0.8879  
Emphysema: AUROC = 0.8263  
Fibrosis: AUROC = 0.7351  
Pleural\_Thickening: AUROC = 0.7992

Hernia: AUROC = 0.9563  
New best model with AUROC: 0.8142  
Epoch 11/20 - Train Loss: 0.3384, Valid Loss: 0.2555  
Accuracy: 0.8982, Mean AUROC: 0.8179  
New best model with AUROC: 0.8179  
Epoch 12/20 - Train Loss: 0.3403, Valid Loss: 0.2540  
Accuracy: 0.8975, Mean AUROC: 0.8178  
Epoch 13/20 - Train Loss: 0.3341, Valid Loss: 0.2547  
Accuracy: 0.8990, Mean AUROC: 0.8180  
New best model with AUROC: 0.8180  
Epoch 14/20 - Train Loss: 0.3320, Valid Loss: 0.2475  
Accuracy: 0.9021, Mean AUROC: 0.8159  
Epoch 15/20 - Train Loss: 0.3322, Valid Loss: 0.2519  
Accuracy: 0.8995, Mean AUROC: 0.8132  
Atelectasis: AUROC = 0.8331  
Cardiomegaly: AUROC = 0.9363  
Effusion: AUROC = 0.8172  
Infiltration: AUROC = 0.6968  
Mass: AUROC = 0.8806  
Nodule: AUROC = 0.8314  
Pneumonia: AUROC = 0.5578  
Pneumothorax: AUROC = 0.8446  
Consolidation: AUROC = 0.7488  
Edema: AUROC = 0.8829  
Emphysema: AUROC = 0.8531  
Fibrosis: AUROC = 0.7378  
Pleural\_Thickening: AUROC = 0.7905  
Hernia: AUROC = 0.9738  
Epoch 16/20 - Train Loss: 0.3263, Valid Loss: 0.2501  
Accuracy: 0.8999, Mean AUROC: 0.8196  
New best model with AUROC: 0.8196  
Epoch 17/20 - Train Loss: 0.3270, Valid Loss: 0.2522  
Accuracy: 0.8982, Mean AUROC: 0.8208  
New best model with AUROC: 0.8208  
Epoch 18/20 - Train Loss: 0.3247, Valid Loss: 0.2527  
Accuracy: 0.8973, Mean AUROC: 0.8229  
New best model with AUROC: 0.8229  
Epoch 19/20 - Train Loss: 0.3204, Valid Loss: 0.2511  
Accuracy: 0.8992, Mean AUROC: 0.8203  
Epoch 20/20 - Train Loss: 0.3200, Valid Loss: 0.2528  
Accuracy: 0.8985, Mean AUROC: 0.8109  
Atelectasis: AUROC = 0.8250  
Cardiomegaly: AUROC = 0.8983  
Effusion: AUROC = 0.8097  
Infiltration: AUROC = 0.6734  
Mass: AUROC = 0.8837  
Nodule: AUROC = 0.8368  
Pneumonia: AUROC = 0.5605  
Pneumothorax: AUROC = 0.8420  
Consolidation: AUROC = 0.7524  
Edema: AUROC = 0.8813  
Emphysema: AUROC = 0.8693  
Fibrosis: AUROC = 0.7317  
Pleural\_Thickening: AUROC = 0.8127  
Hernia: AUROC = 0.9755  
Loaded best model with AUROC: 0.8229

```
Out[16]: {'train_loss': [0.512840187379292,  
    0.3846053685460772,  
    0.3778936295849936,  
    0.3703442132472992,  
    0.36047095111438204,  
    0.355126211302621,  
    0.3526099014282227,  
    0.3487398329802922,  
    0.34231844799859185,  
    0.34071678485189166,  
    0.33843190312385557,  
    0.34026975512504576,  
    0.33408795050212314,  
    0.33195429291043965,  
    0.33215225066457477,  
    0.32628397720200675,  
    0.3270384621620178,  
    0.3247446628979274,  
    0.3204167122500283,  
    0.3200161421298981],  
  'valid_loss': [0.3110121432940165,  
    0.3070401926835378,  
    0.2932629346847534,  
    0.284041223526001,  
    0.27637522260348,  
    0.2689167575041453,  
    0.2626188385486603,  
    0.2585227874914805,  
    0.2542443307240804,  
    0.25621569991111753,  
    0.25548014442125955,  
    0.2539857657750448,  
    0.254738495349884,  
    0.24748066345850628,  
    0.25194777647654215,  
    0.2500927368799845,  
    0.25216431975364684,  
    0.2527491796016693,  
    0.25113846600055695,  
    0.2527521518866221],  
  'valid_accuracy': [0.8932142857142857,  
    0.8932142857142857,  
    0.8932142857142857,  
    0.895952380952381,  
    0.8963095238095238,  
    0.8979761904761905,  
    0.8961904761904762,  
    0.8996428571428572,  
    0.8988095238095238,  
    0.8969047619047619,  
    0.8982142857142857,  
    0.8975,  
    0.8990476190476191,  
    0.9021428571428571,  
    0.8995238095238095,  
    0.8998809523809523,
```

```
0.8982142857142857,  
0.8972619047619048,  
0.8991666666666667,  
0.898452380952381],  
'best_auroc': np.float64(0.8229393220565882)}
```